

PremierPredict-AI: A Stacking Ensemble Approach for Predicting English Premier League Match Outcomes

Chaiyaporn Homtean 67130700346 , Krittin Chaisuvirat 67130700357, and Pawornwit Maneenet 67130700361

SEA612 Artificial Intelligence Fundamentals

Supervised by: Punnarust Silparattanawong

29 March 2026

Abstract—Predicting the outcome of professional football matches remains a significant challenge in sports analytics due to the high variability and multi-factorial nature of the game. Traditional statistical models often struggle to capture the complex non-linear relationships between team form, player market values, and underlying strength metrics. This paper presents PremierPredict-AI, a machine learning system utilizing a Stacking Ensemble architecture to predict English Premier League (EPL) match results (Home Win, Draw, Away Win). By integrating features such as ELO ratings, Transfermarkt market values, and historical performance windows, the model achieves a classification accuracy of 55.3%. Our results demonstrate that the ensemble approach effectively balances variance and bias, surpassing common human benchmarks—including the FPL “Wisdom of Crowds” (54.7%) and BBC expert panels (47.1%)—providing a robust framework for automated sports forecasting.

Index Terms—Machine Learning, Stacking Ensemble, Sports Analytics, Premier League Prediction, ELO Ratings, Meta-Learning, Predictive Modeling.

I. INTRODUCTION

The Global Context (World): The English Premier League (EPL) stands as one of the most unpredictable and data-rich sporting environments globally. Every matchweek provides a wealth of data points; however, predicting outcomes is complicated by “noisy” variables like tactical shifts, red cards, and intense psychological pressure.

Literature Review (Village): Historically, sports forecasting has relied heavily on the ELO Rating system (developed by Arpad Elo) to determine relative skill. While ELO captures historical strength effectively, it lacks the ability to account for sudden squad improvements or financial investments. Recent studies by Peeters emphasize the impact of Market Value (e.g., Transfermarkt data) on long-term league success. Furthermore, the “Wisdom of Crowds” principle suggests that the aggregated opinion of millions of fans (such as Fantasy Premier League users) often provides a reliable baseline for confidence.

Our Approach (Us): This paper documents the design and development of PremierPredict-AI. Unlike simple statistical models, our approach utilizes a Stacking Ensemble architecture, where a team of specialized AI models works together to produce a final consensus. We aim to answer three primary Research Questions (RQs):

- 1) Does a Stacking Ensemble model yield higher accuracy than single models for 3-way match predictions?
- 2) Which deep factors (ELO, form, or market value) exert the highest impact on predictive accuracy?
- 3) How do AI predictions compare against human benchmarks?

II. SYSTEM ARCHITECTURE AND METHODOLOGY

A. Data Pipeline and Featured Variables

To ensure scalability, we implemented a structured Data Pipeline encompassing Data Ingestion, Processing, an Ensemble Engine, and a Dashboard Layer.




Fig. 1. High-level Technical System Architecture of PremierPredict-AI.

We identified four “pillars” of predictive power:

- **ELO Differential:** The relative gap in skill ratings (History).
- **Market Potential:** The financial valuation of the squad (Potential).
- **Momentum (L5 Form):** Averaged points and goal differences over the previous five matches.
- **Crowd Sentiment:** Aggregated FPL Selection and Captaincy percentages.

Theoretical Justification: ELO tells us how a team has performed, while Market Value tells us how good their players inherently are. When these two indicators diverge, it serves as a primary mathematical signal for identifying potential upsets or “dark horses.”

B. Development Roadmap (V1.0 - V5.0)

The model underwent four major tuning phases:

- 1) **V1.0 Baseline (44.8%)**: Used simple Logistic Regression; suffered from high variance and struggled with away matches.
- 2) **V2.5 Feature Engineering (51.2%)**: Integrated Transfermarkt data; surpassed the BBC Experts (47.1%).
- 3) **V3.8 Advanced Single Model (53.9%)**: Switched to XGBoost; robust but still trailed the FPL Crowd (54.7%).
- 4) **V5.0 Stacking Ensemble (55.3%)**: By combining models, we captured nuanced patterns that single models missed.

C. The Stacking Architecture

This approach treats machine learning as a hierarchical "team effort":

- **Layer 1 (Base Learners)**: We selected models with diverse mathematical foundations. Random Forest uses bagging to provide stable, low-variance predictions. XGBoost efficiently handles missing data and finds complex non-linear patterns. Gradient Boosting (GBDT) strictly focuses on correcting the errors of preceding models.
- **Layer 2 (The Meta-Model)**: A Logistic Regression "Manager" evaluates the predictions of Layer 1. It learns contextual trust (e.g., favoring XGBoost when its probability threshold exceeds 80%), providing the final 1.4% accuracy boost.

III. EXPERIMENTAL RESULTS

Evaluation was conducted on a testing set of 380 matches from the 2024 EPL season.

TABLE I
PERFORMANCE COMPARISON

Forecaster	Avg Accuracy	Macro F1
Chris Sutton (BBC Guru)	47.1%	0.48
FPL "Wisdom of Crowds"	54.7%	0.52
Single XGBoost Model	53.9%	0.54
PremierPredict-AI V5.0	55.3%	0.56

To understand feature sensitivity, we calculated the correlation between individual domains and the final match outcome.

TABLE II
CORRELATION TO OUTCOME

Feature Domain	Correlation	Description
ELO Diff	0.61 (High)	Relative Skill Gap
Market Value	0.42 (Med)	Depth of Talent
Home Field	0.32 (Med)	Environmental Factor
FPL Captaincy	0.27 (Low)	Social Trust

IV. DISCUSSION & CHALLENGES

Despite reaching 55.3% accuracy, we encountered several bottlenecks that prevented reaching the 60% threshold:

A. Analysis of "Wisdom of Crowds"

A key finding was the difficulty of beating human fans (54.7%). Humans are surprisingly adept at noticing contextual nuances, such as fatigue following a mid-week game. However, our AI surpassed the crowd by 0.5% primarily by remaining emotionally neutral during big "Derby" matches, where fan bias frequently clouds judgment.

B. The Draw Dilemma

Predicting a "Draw" is the graveyard of football models. Our model's precision for draws (51.2%) is significantly lower than for Wins (71.3%). This confirms that draws are disproportionately influenced by randomness, failure to convert chances, or mid-game red cards, which historical data cannot anticipate. This discrepancy highlights that draws are often the result of stochastic match events rather than long-term team strength.

C. Probability Calibration and the "Big Six" Bias

In early versions (V1.0), the model exhibited extreme bias towards "Big Six" teams like Manchester City, predicting their victory even during slumps. By applying Probability Calibration (as seen in the logic flow in Fig. 2), the Meta-Learner improved the Log-Loss metric to 0.82 (down from the base learners' average of 0.93). This made the AI more "honest" about its confidence levels, significantly reducing false-positive predictions for elite clubs.




Fig. 2. End-to-End Prediction Logic Flowchart.

D. Real-time Impact Limitation

The current architecture relies on data available 24 hours before kickoff. It operates completely blind to the official starting lineups. Integrating a Live API for "Team Sheets" (announced 1 hour before kickoff) is projected to provide an additional 2-3% accuracy gain.

V. CONCLUSION AND FUTURE WORK

The PremierPredict-AI project successfully demonstrates that a Stacking Ensemble of machine learning models can outperform traditional expert analysis and aggregated human sentiment. By bridging the gap between historical ELO ratings and financial market valuations, we created a model that generalizes effectively across the chaotic environment of the English Premier League. While the "Beautiful Game" will always retain an element of surprise, our results show that 55% of that surprise is a predictable pattern hidden within the data.

Future Work:

- **Deep Learning (LSTM):** To capture the sequential nature of a season, where confidence builds over multiple weeks.
- **Explainable AI (XAI):** Implementing SHAP or LIME visualizations so users can see exactly why a team like Chelsea is predicted to lose despite a high market value.
- **Social Sentiment:** Analyzing Twitter/X data to capture fan anxiety or excitement as a predictive feature.

ACKNOWLEDGMENTS

The authors would like to thank Punnarust Silparattanawong for providing the theoretical foundation in the SEA612 Artificial Intelligence Fundamentals course.

REFERENCES

- [1] D. Berrar, "Tree-based Ensembles for Sports Analytics," *Journal of Prediction Markets*, 2022.
- [2] T. Peeters, "Estimating the Economic Value of EPL Squads," *Journal of Sports Economics*, 2023.
- [3] Football-Data.co.uk, "EPL Match Statistics Database," 2024. [Online]. Available: <http://www.football-data.co.uk>. [Accessed: March 15, 2026].
- [4] M. J. Dixon and S. G. Coles, "Modelling Football Scores," *Applied Statistics*, 1997.
- [5] ClubELO, "The Global Club ELO Index," 2024. [Online]. Available: <http://clubelo.com>. [Accessed: March 20, 2026].
- [6] Transfermarkt, "Squad Market Value Report," 2024. [Industry Data]. Available: <https://www.transfermarkt.com>. [Accessed: March 22, 2026].